ShannonProver: Towards Automating Formal Cryptographic Proofs

Yiping Ma[†], Yu-Lin Tsai[†], Mayank Rathee[†], Deevashwer Rathee[†],
François Dupressoir[‡], Pierre-Yves Strub[§], Raluca Ada Popa[†]
[†]*UC Berkeley* [‡]*University of Bristol* [§]*PQ Shield*

arXiv:2607.02847v4 [cs.CR] 8 Oct 2026

***Abstract*—Cryptographic proofs are produced at a scale that increasingly exceeds the community's ability to verify them manually. Machine-checked proofs offer a path toward scalable proof verification, but they shift the bottleneck to writing the proofs themselves: even when the high-level proof plan is known, turning it into a proof script requires spelling out every detail the plan leaves implicit, which is laborious even for experts. This paper presents ShannonProver, an agentic system for automating cryptographic proofs. Given a protocol and its security definitions modeled by a cryptographer, ShannonProver decomposes the main theorem into intermediate games and lemmas, and constructs EasyCrypt proof scripts for those lemmas.**

We evaluate ShannonProver on a new dataset of lemmas in EasyCrypt. The benchmark spans textbook primitives, deployed standardized protocols, as well as recent NIST proposals, and includes expert case studies drawn from a corpus that has not previously been available online. On case studies such as ChaChaPoly1305 and MEE-CBC, ShannonProver completes within hours proof developments that historically took experts weeks to months. More broadly, this work suggests a path toward accelerating cryptographic research: as agents automate the proof-engineering burden, cryptographers can iterate more quickly on new constructions, obtain machine-checked assurance earlier, and bring protocols from design to deployment faster.

1. Introduction

Cryptographers *"generate more proofs than we carefully verify"*, as Halevi observed two decades ago [1]. This concern is even more pressing today. Modern cryptographic protocols are increasingly complex, their proofs are longer and more intricate, and their correctness has consequences far beyond research papers: they have become standards [2–5], libraries [6–9], and security infrastructure [5,10–12]. Flaws discovered after publication or even standardization underscore the same point [13,14]: the community cannot manually scrutinize cryptographic proofs at the same scale at which it produces them.

Machine-checked proofs offer a way to scale proof verification. Here, a proof is written as a *proof script*, a sequence of commands called *tactics*, and a proof assistant [15–17] checks the script mechanically. This work uses EasyCrypt [18], a proof assistant designed for game-based cryptographic proofs. If the check passes, the theorem that the script proves is certified without any expert examining the argument by hand. However, the bottleneck shifts to producing the proof script in the first place.

Writing such proof scripts is far more exacting than writing traditional pen-and-paper proofs: a machine-checked proof is not a line-by-line translation of a pen-and-paper proof; rather, it expands the cryptographer's shorthand of a cryptographic argument (e.g., "two games are indistinguishable") into explicit stateful programs and precisely reasons about the program equivalence. Historically, developing formal proofs for standardized schemes such as MEE-CBC [19], CMAC [20], and ChaCha20-Poly1305 [21] has required several person-weeks to several person-months[1], and the effort for recent post-quantum NIST standards such as ML-KEM [13,22] is even larger, involving more than 10 experts for one or two years.

This effort raises a natural question: can experts delegate writing these proof scripts to LLM agents? Machine-checked proofs make such delegation safe: the agent performs the proof search, which may be heuristic, while EasyCrypt performs the proof checking, which is sound. The agent may hallucinate or make mistakes along the way, but this cannot compromise soundness: EasyCrypt rejects every invalid proof attempt, so once the agent produces a script that passes the check, the theorem is certified.

The most direct approach to this question is to ask the agent to write the entire proof through checker-in-the-loop interaction [23,24]: the agent receives the specification of the scheme and the statement of the main theorem, submits tactics to EasyCrypt, receives the resulting goal or error, and repeats until the theorem is proved. While this approach suffices for simple textbook schemes (e.g., ElGamal), it falls far short of our target: formal proofs at the scale and complexity of NIST standards and proposals. We observed that even strong LLMs fail on such proofs. They sometimes re-prove facts they have already established or drift to the wrong level of abstraction; as a result, their proofs grow far longer than necessary, or they eventually give up.

Experts do not develop such proofs monolithically. They proceed in three phases (Fig. 1): they first write the *specification*, which formalizes the scheme, its security definitions,

1. The exact historical effort for the ChaChaPoly development is no longer recoverable; the person-month estimate is based on the size and complexity of the current proof relative to comparable projects with known effort.

and the underlying assumptions (Phase I), then decompose the main theorem into intermediate games and lemmas (Phase II), and finally prove each lemma with a tactic-level proof script (Phase III). Phase II requires cryptographic insight and high-level planning, whereas Phase III is tedious and demands exacting, mechanical proof writing. The two phases are not sequential: whether a decomposition works is often known only after trying to prove its lemmas. Experts therefore move back and forth between the two phases until every lemma is proved. This process is costly: the original MEE-CBC proof took experts a few person-weeks, and the ChaCha20-Poly1305 proof a few person-months.

In this work, we develop ShannonProver, an agentic system that follows this workflow to automate Phases II and III. It has two components, both built on LLM agents (Fig. 2): a *decomposer* for Phase II and a *lemma prover* for Phase III. The decomposer works at the level of the whole proof: it writes the intermediate games and lemma statements, and hands individual lemmas to the lemma prover. The lemma prover proves each lemma by interacting with EasyCrypt tactic by tactic. If it succeeds, it returns a verified proof script; otherwise, it returns feedback, which the decomposer uses to revise or split the lemma, or to restructure the decomposition. Because the decomposer can launch the lemma prover on several independent lemmas at once, ShannonProver greatly reduces the time needed to test whether a decomposition actually works. Given only the specification and the main theorem, ShannonProver constructs complete proofs for MEE-CBC in 2 hours and for ChaCha20-Poly1305 in about 5 hours.

This workflow hinges on the lemma prover: a decomposition can be tested only as quickly and as reliably as its lemmas are proved. Yet even on a single lemma in Phase III, the checker-in-the-loop approach can fail when the lemma is difficult or tedious. These failures have two sources: the agent lacks the context relevant to each step of the proof, and it explores the proof along a single linear path. The lemma prover addresses them with two design ideas: a *proof-state compiler* and *multi-agent tree search*.

❶ Proof-state compiler (§3). The context an agent needs changes with every tactic. Each tactic transforms the proof state, and with it which *proof resources* apply (previously proved lemmas, assumptions, and hypotheses; App. A.1) and how their arguments must be instantiated (left in Fig. 4). An agent can recover this context only by rereading the project files after every tactic, which is slow; when it misses something, it often switches to a worse strategy or gets stuck. Yet this context is fully determined by the current state and the project: the agent need not rediscover it, because it can be computed. Much as a compiler resolves names and types so that programmers need not, the *proof-state compiler* (right in Fig. 4) analyzes each new state and shows the agent the resources that apply, with their arguments instantiated. The agent keeps every strategic decision; the compiler takes over the bookkeeping.

❷ Multi-agent tree search (§4). A proof rarely proceeds in a straight line: a step that looks reasonable may reveal

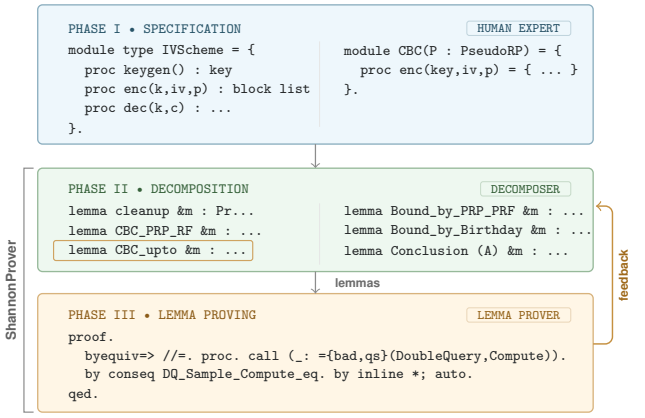

Figure 1. The three phases of a formal proof development. A human expert writes the specification (Phase I), and ShannonProver automates decomposition (Phase II) and lemma proving (Phase III).

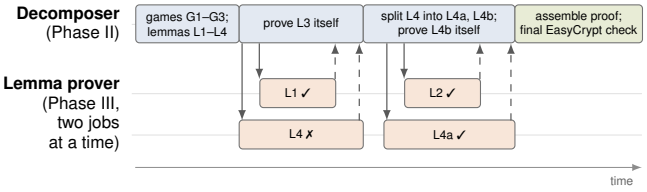

Figure 2. Interleaving of Phases II and III in ShannonProver. Solid arrows hand a lemma to the lemma prover; dashed arrows return the result, either a proof (✓) or feedback (✗). The decomposer does not wait for results: while L1 and L4 are being attempted, it proves L3 itself. When the attempt on L4 fails, the decomposer splits L4 into L4a and L4b. Two jobs run at a time here for illustration; the number of concurrent jobs is configurable.

only much later that it has led to a dead end. A single agent explores one path at a time, and it often fails to notice when to turn back: it keeps trying variants of the same failing move instead of returning to where its strategy went wrong. The lemma prover therefore organizes the search as a tree of checker-accepted proof states (Fig. 6). Agents can backtrack within their own branches, and a deterministic *orchestrator* oversees all branches: when an agent stalls, it starts a fresh agent on a different branch, prunes branches that no longer make progress, and records dead ends so that no agent retries a failed approach at the same state. Each agent reasons about its own branch; the orchestrator decides where the search should go.

We view this work as a first step toward automating formal proofs for cryptography. Ultimately, we hope ShannonProver will let experts focus on the critical judgments in formal verification, while automating as much of the remaining work as possible.

2. Design Overview

2.1. Observations in Agentic Proving

We observe the failures below in baseline runs, where the agent interacts directly with EasyCrypt (left in Fig. 4).

They fall into the two groups identified in §1: the agent lacks the context relevant to each step of the proof, and it explores the proof along a single linear path.

Missing context. At every step, the agent must decide which proof resources apply, which one to use, and how to apply it. Making this decision requires assembling information scattered across project lemmas, imported theories, module arguments, and even the source code of the proof checker. When the agent lacks this information, we observe three recurring failure patterns, each at a different point of this decision: the agent misses a resource that applies, chooses one that does not, or cannot apply the one it has chosen.

❶ *Unconscious lowering.* A proof often moves across several abstraction levels (Fig. 3): from probability algebra and pRHL equivalences, to module or oracle calls, procedure bodies, and finally SMT side conditions. Higher levels let the prover reuse facts already proved about whole games, modules, or oracle behaviors. Human experts descend to a lower level only when the proof really needs more concrete program structure, guided by experience and by the intuition they gained from lemma decomposition (Phase II); an LLM agent has to make this judgment after every tactic. The agent can mistakenly interpret a local but unhelpful goal change as real proof progress. In the baseline runs, a tactic such as `inline*` is very tempting because it is easy to write and immediately produces a different-looking goal. What the agent does not see is that this change may cut off resources that are usable at the higher level: a lemma stated about a procedure call, such as an equivalence between two oracle implementations, applies only while that call is still visible in the goal; once the call is inlined, the lemma no longer matches (App. D). The agent must then either reprove from scratch what the lemma would have established, or keep pursuing the original plan that relied on the lemma but can no longer succeed from this state. The failure is not that the agent chose to descend; descent is often necessary. The failure is that the interface gives no state-aware information when a locally productive-looking move is globally destructive.

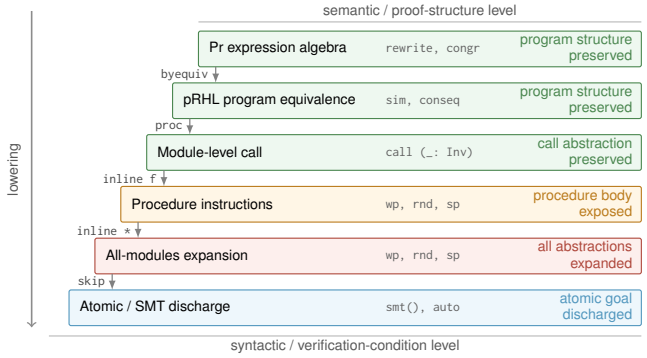

Figure 3. The tactic pyramid: abstraction levels in an EasyCrypt proof, from probability expressions down to SMT goals. Each level lists the tactics used within it; each arrow between two levels is labeled with the tactic that moves the proof down. Lowering exposes program detail but can invalidate resources that apply only at higher levels.

❷ *Semantic mislocalization.* An agent can choose a resource that does not apply. This is not simply a lack of information: the proof context may contain the right clues, but the agent can be pulled toward the wrong resource by names that appear in the goal or by the literal wording of an error message. For example, in one run, a goal mentioning `IFinRO` and `IndRO` led the agent to treat these names as evidence for an eager/lazy random-oracle lemma and to spend the rest of the run trying to fit `pr_RO_FinRO_D`, which is unrelated to the proof (App. D).

❸ *Agent flinch.* An agent can choose the right resource and still fail to apply it. Applying a lemma means turning the choice into a command that EasyCrypt accepts: looking up the lemma's exact statement, supplying its module arguments, matching types, and finding the right tactic form. When this information is available only through trial and error with the checker, the agent often abandons the plan after a few rejected attempts and switches to a tactic that is easier to write, even if it leads to a worse proof state (App. D). The failure is not the absence of a proof idea, but the cost of carrying it out.

Exploring a single path. Even with the right context at every step, a tactic that EasyCrypt accepts can lead the proof into a dead end that becomes visible only many steps later. The agent must therefore also decide when to abandon its current strategy and where to resume. Even an agent that can return to an earlier state often does not notice when it should: it stays at a state and keeps submitting variants of the same failing move, or keeps repairing a branch whose strategy no longer works, instead of returning to the point where the strategy went wrong. And when it does backtrack, it pursues one strategy at a time, so several plausible strategies can only be tried one after another. These failures do not come from a single wrong step, but from the lack of any view of the search as a whole.

2.2. Design Ideas

Proof-state compiler. The three failure patterns point to an interface problem: the agent must reconstruct, through ad hoc queries such as `Grep` and `Read`, proof context that is already determined by the current proof state and the project. ShannonProver instead computes this context deterministically with a *proof-state compiler* (§3, Fig. 4). At each proof state, the compiler takes as input the checker's output, the proof history, and the project context, including source files, imported theories, module definitions, and available lemmas. Its output addresses the three failure patterns in turn: it shows which resources still apply at the current state, so that the agent does not lower past them; it matches resources against the goal by structure rather than by name; and once the agent selects a resource, it fills in the arguments and premises that the current state determines, so that the agent can carry its plan through. These outputs are facts about the current proof state, not a proof strategy: the agent still decides which proof operation and resource to use next. The compiler's analyses follow classical compiler design,

which builds an intermediate representation, resolves names and types, tracks which facts remain valid, and reports diagnostics [25]; Fig. 5 shows its three layers.

Multi-agent tree search. Neither failure of single-path exploration can be fixed inside one agent: an agent that does not notice when to turn back cannot be relied on to turn back, and a single agent can pursue only one strategy at a time. ShannonProver therefore organizes the search as a tree of checker-accepted proof states (§4, Fig. 6). Each agent can backtrack within its own branch, which handles the wrong turns it recognizes. A deterministic orchestrator watches the progress of all branches and handles the rest: when an agent stalls, it starts a fresh agent on a different branch, and because several branches run at once, plausible strategies are tried in parallel rather than one after another. It also prunes branches that fall behind and records the attempts that failed.

3. Proof-State Compiler

Writing an EasyCrypt proof requires two kinds of work. First, the prover makes a strategic choice: it selects a proof operation and the proof resources with which to carry out that operation. By a *proof operation* we mean a kind of proof step rather than one specific command: for example, using a lemma about a procedure call (`call`, `ecall`), rewriting with an equation (`rewrite`), or turning a probability statement into a program equivalence (`byequiv`). The same operation can be written with different tactics and tactic forms. A *proof resource* is something the proof can use instead of reasoning from scratch, such as a lemma already proved in the project or in a library, an axiom of the specification, or a hypothesis in the current proof. Which resources can be used changes as the proof proceeds (App. A.1). Second, the prover must realize that choice as an EasyCrypt command that is valid at the exact current proof state. This requires resolving names and modules, supplying the appropriate arguments, and using the correct tactic form.

The first kind of work is *strategic*: it requires proof insight, that is, semantic reasoning about what the programs and statements mean. The second is *mechanical*: it is often determined by the syntactic form of the current goal and of the selected resource, and therefore changes with every proof state. Repeatedly working on such details through trial and error consumes reasoning budget on work that can instead be performed deterministically. More importantly, we observed that these low-level failures can cause an agent to abandon an otherwise correct high-level plan and fall back to a worse proof strategy, as in the agent flinch of Section 2.

We therefore separate strategic choices from their mechanical realization. Concretely, the agent chooses the proof operation, the proof resource, and any other arguments that require proof insight, such as an invariant. The proof-state compiler then *realizes* the agent's intended step: it keeps the proof operation and resource that the agent chose, fills in what the current state determines, and uses the tactic form that EasyCrypt expects; everything else is left to the agent. In particular, the compiler never replaces the agent's choices. Changing the proof strategy is also the agent's decision; Section 4 describes the backtracking that supports it.

Running example. We use a simple losslessness proof throughout this section to illustrate our design (Fig. 4). The goal is to prove that the procedure `D2(O).distinguish`, shown below in simplified form, is lossless:

```
proc D2(O).distinguish() = {
  D2(O).O.init();
  b <@ Adv_MAC_to_F(A, D2(O).O).guess();
  return b;
}
```

Elsewhere in the project, the following lemma has been proved:

```
lemma Alossless_F (X <: PRF_Oracles) :
  islossless X.f =>
  islossless Adv_MAC_to_F(A, X).guess.
```

After opening the procedure with `proc`, the agent correctly decides to use the proof operation `call` with the resource `Alossless_F`, but submits `call (Alossless_F _ _)`. EasyCrypt rejects the command because it cannot infer the module argument. The intended step should be realized as

```
call (Alossless_F (<: D2(O).O) _).
```

In the agent's attempt, the first hole is a module argument, which the current state determines. The second hole is a premise of the lemma. The compiler fills it only if the current proof context already has a hypothesis stating that premise; otherwise the premise is a new proof obligation that requires semantic reasoning, and it is left to the agent.

Design overview. Our compiler is organized into three layers (Fig. 5). The first layer, state projection (§3.1), determines which proof state the agent is working on; all later analysis refers to this state. The second layer, ProofIR (§3.2), parses the current goal into a typed representation, together with the proof operation and resource of the agent's attempt, if there is one. The third layer, resource and binding analysis, uses these results in one of two ways, depending on whether the agent has already selected a resource. If it has not, *resource exposure* (§3.3) finds the resources that apply at the current state and shows them to the agent. If it has, *binding analysis* (§3.4) realizes the agent's intended step.

3.1. State Projection: Where am I in the proof?

Writing a formal proof for a lemma is often highly local. At each step, the information needed to construct the next tactic sequence is often specific to the current proof state. An accepted tactic advances the proof to a new state, a rejected tactic reports an error while leaving the state unchanged, and an undo restores an earlier state. Consequently, the interaction transcript between the agent and EasyCrypt interleaves goals, error messages, and analyses associated with different proof states.

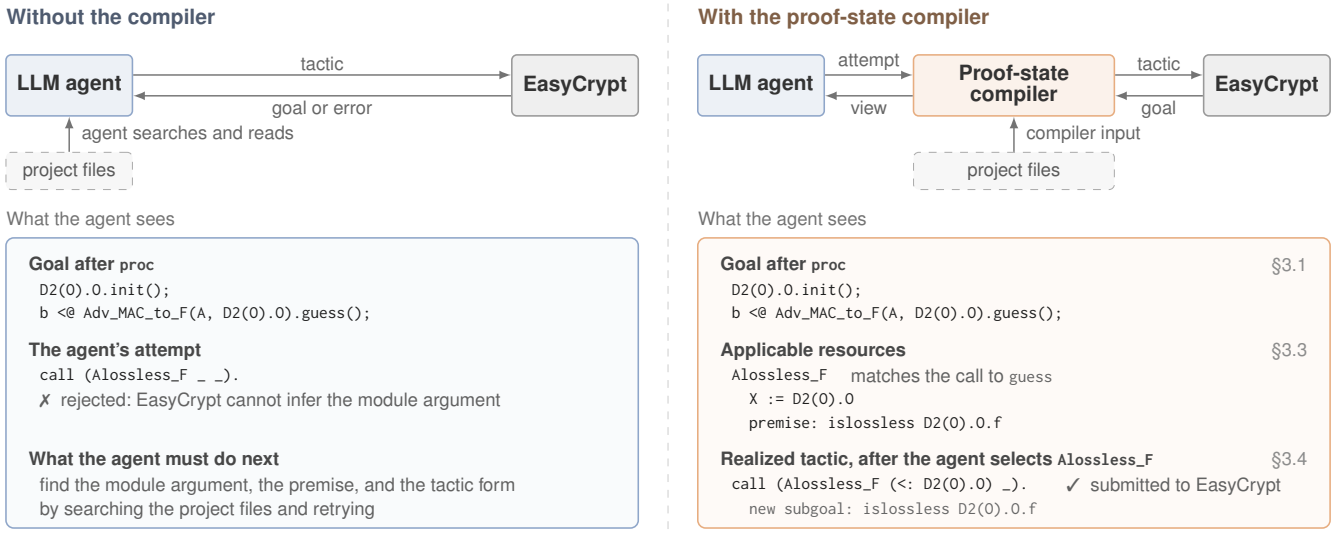

Figure 4. One step of the running example without (left) and with (right) the proof-state compiler. Left: the agent's attempt `call (Alossless_F _ _).` is rejected because EasyCrypt cannot infer the module argument, and the agent must search the project files to recover it. Right: the compiler shows the resources that apply at the current state, with their arguments instantiated (§3.3); once the agent selects `Alossless_F`, the compiler realizes the tactic and submits it (§3.4). The agent chooses the resource; the compiler fills in only what the state determines and leaves the lemma's premise to the agent as a new subgoal.

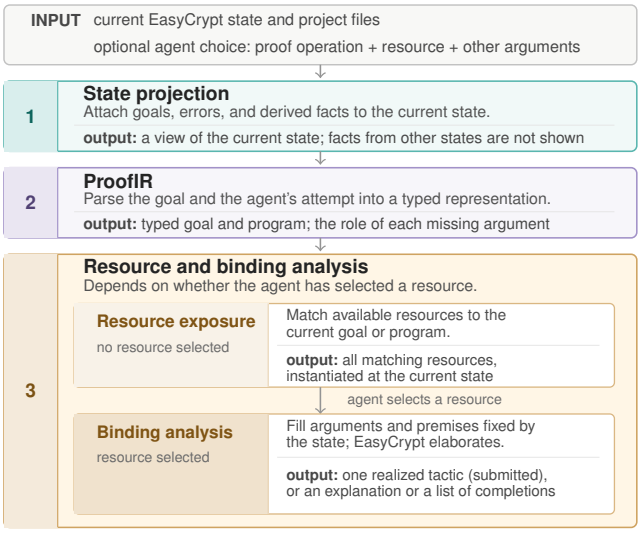

Figure 5. The three compiler layers. The third layer either shows the agent the resources that match the current state or realizes a resource that the agent has already selected.

We do not ask the agent to keep the history. Instead, the compiler maintains a state-centered view consisting of the current goal together with information derived specifically for that state. When the proof advances or returns to an earlier state, the view is recomputed for that state.

This design serves two purposes. First, it removes bookkeeping from the agent: the agent does not need to determine which parts of a long transcript are still relevant. Second, it prevents stale information from silently influencing later proof decisions. An error produced before an undo, for example, should not become a general statement that a tactic or resource is unusable.

Long-term reasoning is not the compiler's job. The agent keeps its overall plan and the proof strategy it is following, and the proof tree of Section 4 records which alternatives have already failed. State projection only ensures that the proof facts presented by the compiler describe where the agent is *now*.

Running example. After the agent runs `proc`, the goal shows the body of `D2(O).distinguish`. The rejected attempt `call (Alossless_F _ _).` does not change the proof state, so its error is attached to the state after `proc`, and the compiler analyzes that same state. If the agent later undoes `proc` or moves to a different state, this error is no longer shown as if it described the new state.

3.2. ProofIR: How to interpret the state?

Once the current state is fixed, the compiler must determine what part of the agent's proof attempt is already a strategic choice and what part is an incomplete mechanical detail. Raw tactic syntax does not make this distinction explicit. An underscore in an EasyCrypt term, for example, may stand for a module parameter, a value argument, or an unresolved proof premise. Treating all of these simply as "missing syntax" would give the compiler no principled way to decide what it may fill in.

ProofIR makes these roles explicit. For each proof attempt of the agent, it records three kinds of information:
❶ The current goal and relevant program structure.
❷ The proof operation and proof resource selected by the agent.

❸ The typed parameters of that resource, including which arguments were supplied and which remain unresolved.

If the agent has not submitted an attempt, ProofIR records only the first kind of information, which resource exposure (§3.3) uses.

ProofIR turns an informal failure such as "this command did not elaborate" into a precise state-dependent question about the proof operation and resource that the agent has already selected.

Running example. The ProofIR layer parses the agent's attempt `call (Alossless_F _ _)`. It records that the requested proof operation is `call` and the selected resource is a lemma `Alossless_F`. It parses the lemma as

```
module param  : X
module type   : PRF_Oracles
premise       : islossless X.f
conclusion    : islossless Adv_MAC_to_F(A, X).guess
```

In this representation, the role of each missing argument in the agent's attempt is explicit: the first hole asks for a module instantiation; the second asks for a proof of a losslessness premise. This layer only parses: it does not fill either hole or assert that the lemma applies. It makes explicit what binding analysis (§3.4) must relate.

3.3. Resource Exposure: What can I use?

Resource exposure handles the case in which the agent has not yet selected a resource. This addresses unconscious lowering and semantic mislocalization (§2), in which the agent misses a resource that applies or chooses one that does not. For example, after a procedure body is opened, the goal shows a procedure call, but it does not show which existing lemmas can be used at that call.

The compiler therefore finds resources that are *structurally* applicable. This matching is syntactic: the compiler compares the form of the current goal with the form of each resource's statement, rather than by the similarity of names or text. Because the matching is syntactic, we do it deterministically instead of using an LLM. It proceeds in three steps.

❶ The compiler finds the objects in the current goal against which resources can be matched, such as procedure calls. In the running example, after the proof operation `proc` and before the agent submits any `call`, the goal contains the two statements of the procedure body shown above. The last statement is the procedure call `Adv_MAC_to_F(A, D2(O).O).guess`, which is where a lemma about a call can be applied with `call`.

❷ The compiler finds the resources available in the current file and its imports (App. A.1) whose statements match that call:

```
resource     :   Adv_MAC_to_F(A, X       ).guess
current call :   Adv_MAC_to_F(A, D2(O).O ).guess
```

It checks that the procedure names are the same (`guess`), that the enclosing module application is the same (`Adv_MAC_to_F`), that the module arguments already present are the same (`A`), and that the remaining module parameter `X` can be instantiated with the concrete module `D2(O).O` from the current state.

❸ The compiler shows the agent each matching resource, with its statement instantiated at the current state, including the premises that remain to be proved. Here, the agent is shown `Alossless_F` with

```
premise   : islossless D2(O).O.f
conclusion: islossless
              Adv_MAC_to_F(A, D2(O).O).guess
```

If several resources match, the compiler shows all of them, without ranking them, and the agent chooses. Once the agent selects a resource, binding analysis (§3.4) completes the tactic.

3.4. Binding Analysis: How should I execute?

Binding analysis handles the case in which the agent has selected a proof operation and a resource, either on its own, as in the running example, or from those shown by resource exposure. It realizes the agent's intended step. This addresses agent flinch (§2): the agent has chosen the right resource but cannot turn the choice into a command that EasyCrypt accepts. We allow the compiler to fill an argument only when that argument is determined by the selected proof operation, the selected resource, and the current proof state. The analysis proceeds in three steps.

❶ It determines where the attempted proof operation applies in the current state and checks whether the selected resource applies there. This applicability is inherently state-dependent: a lemma about a procedure call, for example, can be used only while the corresponding call is still visible in the goal; after the procedure is inlined or the proof moves past that call, the same lemma may no longer apply (the unconscious lowering of §2; Fig. 3).

❷ When the selected resource does apply, the compiler uses its typed structure to determine missing arguments. Module parameters, procedure parameters, or other bindings can often be read off by comparing the resource with the program object to which it is being applied. A premise of the resource is filled only if a hypothesis in the current proof context states exactly that premise. This is mechanical: the agent has chosen the resource, and the hypothesis is determined by the premise. Otherwise the premise is left as a hole, because it is a new proof obligation. Finally, the compiler writes the instantiated resource in the form that the proof operation expects (for example, marking a module argument with `<:`).

❸ The compiler asks EasyCrypt to elaborate the instantiated term. If elaboration fails, the proposed binding is discarded; the compiler does not substitute a different resource. If this yields exactly one tactic, the compiler submits it to EasyCrypt. Otherwise, it tells the agent why no tactic was produced, for example that the selected resource does not apply at the current state; if several completions are possible, it lists them and lets the agent choose.

Running example. As in §3.3, comparing the call `Adv_MAC_to_F(A, D2(O).O).guess` with the statement of the selected lemma determines the module argument uniquely as `X := D2(O).O`, which leaves the premise `islossless D2(O).O.f`. If the current proof context has a hypothesis `Ofll : islossless D2(O).O.f`, for example an assumption of the lemma being proved, the compiler uses it, and the realized tactic is

```
call (Alossless_F (<: D2(O).O) Ofll).
```

Otherwise, the premise is a new proof obligation. The compiler leaves it as a hole, and the realized tactic is

```
call (Alossless_F (<: D2(O).O) _).
```

EasyCrypt then turns the hole into a new subgoal, `islossless D2(O).O.f`, which the agent proves.

4. Tree-Based Proof Orchestration

Section 2 observed that a single agent explores the proof along one path and often does not notice when it should turn back. ShannonProver addresses this with a tree-based search explored at two levels. Within a single agent, the agent can backtrack and re-explore an earlier state when it discovers that the current strategy is unproductive. Across agents, an orchestrator manages several branches of the proof tree: it decides when to start a new agent on a fresh branch, when to stop an unpromising one, and how failures found on one branch should guide the others. We first introduce the basic tree concepts (§4.1).

4.1. Proof Tree

A *node* represents a proof state. Working on a proof means moving progressively between nodes. An *edge* is an accepted transition between states: the agent submits a tactic, the checker accepts it, and the proof state moves from the current node to a child node. A rejected tactic leaves the state unchanged and creates no edge. A *path* is the sequence of nodes from the root to the current node, representing proof attempts. An *agent* is a prover instance that extends a single branch of the tree, advancing the proof one node at a time.

The search space is a tree rather than a single line because the same state can be continued in more than one way. A node gains another child in two ways. First, an agent may *backtrack* to that node and try a different strategy; if the new tactic is accepted, it leads to a different child state (§4.2). Second, the orchestrator may *spawn* another agent at the same shared state, so that a different strategy can be explored without waiting for the current branch to finish (§4.3). The resulting branches are *siblings* below a common node. A branch ends at a *leaf*: either a closed proof (`qed.`) or a state that is abandoned because it is a dead end or is pruned by the orchestrator (§4.3).

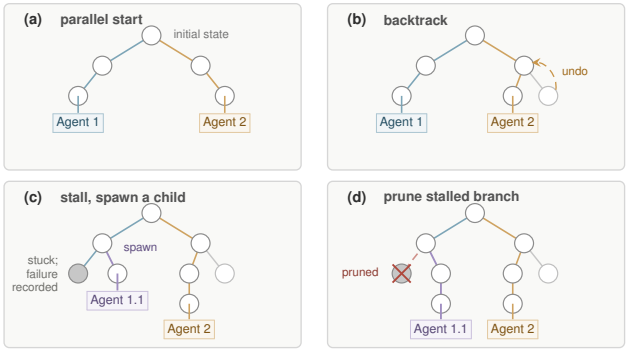

Figure 6. A proof tree evolving over a search. **(a)** Two agents explore in parallel from the initial state. **(b)** Agent 2 recognizes a wrong turn, backtracks to an earlier state, and takes a different edge. **(c)** Agent 1 stalls; the orchestrator records the failed attempt and spawns Agent 1.1 at the parent state. **(d)** The stalled branch is pruned (×); Agent 1.1 and Agent 2 continue.

4.2. Backtracking within an Agent

At each node, the agent reads the output of the proof-state compiler (Section 3), which gives it a structured view of the current proof state. Based on this view, the agent extends its branch one tactic at a time, then reads the compiler's output for the resulting state and continues from there.

An agent does not only move forward. During proof construction, it may discover that the current strategy has become unproductive, or that an earlier tactic quietly led the proof into a dead end. In that case, the agent can backtrack to the earlier state, revise the proof from that point, and extend a new edge to a new child. This makes exploration less brittle: the agent does not have to make a perfect sequence of irreversible choices, but can try a strategy, learn from where it breaks, and return to a better branching point.

4.3. Scheduling across Agents

Backtracking (§4.2) handles the wrong turns that the agent itself recognizes. It does not address the two limitations observed in Section 2: an agent may stall at a state instead of backtracking, and a single branch explores alternatives one after another, even when several strategies are worth exploring together.

To address both limitations, ShannonProver uses an orchestrator. The orchestrator is not an LLM agent, but instead a deterministic algorithm that manages the same proof tree from above. It decides where to fork the search with a fresh agent, which branches should be stopped, and how failures on one branch should inform the others. The orchestrator uses three mechanisms: spawning, pruning, and negative memory (Fig. 6).

Spawning children. At each node, the orchestrator may spawn additional children, each exploring a different path. Without control, this would quickly become infeasible: spawning two children from the root, then two children from

each child, would make the number of live agents grow exponentially. ShannonProver therefore spawns new agents only when the current search behavior suggests that another strategy is worth trying.

❶ *Persistent local failure.* If an agent at a node (state) repeatedly submits tactics and the checker rejects them, this suggests that the agent may be stuck in a local neighborhood of bad moves and the current branch is unproductive. Rather than letting the agent spend more time on the same strategy, the orchestrator creates a sibling branch. It does so by spawning a child *at the parent* of the node where the agent is currently working, giving the search a chance to approach the state through a different path.

❷ *Structural undo.* Some tactics fix the shape of the rest of the proof: they open a procedure or turn a probability statement into a program equivalence (e.g., `proc`, `byequiv`), cut the program with an intermediate assertion or invariant (`seq`, `while`, `call`), inline code (`inline`), or introduce an intermediate fact (`have`, `transitivity`). We call an undo *structural* when it undoes such a tactic: the agent has abandoned an earlier decision that shaped much of the later proof, so the state before that tactic is a valuable branching point. If the same agent has not repaired its strategy within five minutes, the orchestrator spawns a child at that state.

Pruning. Spawning creates new opportunities, but it also creates the risk of an unbounded search. The orchestrator therefore prunes branches and terminates unproductive agents. The pruning decision is based on a *progress gap*: if an agent remains far behind the other agents in the tree, it becomes a candidate for pruning.

This decision has to be made carefully. A branch that looks behind is not necessarily useless: an agent may be backtracking through earlier states, which temporarily reduces its depth but may eventually lead to a better proof strategy. The orchestrator therefore distinguishes agents that are temporarily behind because they are performing meaningful backtracking from agents that remain behind throughout the search. In this way, pruning keeps the number of live agents bounded without discarding every branch that briefly moves backward.

Negative memory. When one agent fails to extend a path, the other agents should not blindly repeat the same experiment. The orchestrator therefore records failures as state-indexed *negative memory*. For a given node, negative memory records which attempted extension led to a dead end and why. When the orchestrator later spawns a child at that node, it passes the corresponding negative memory to the new agent, so that the new branch can avoid spending budget on the same failed strategy.

Negative memory is a search-control mechanism, not proof knowledge. It records that a particular experiment failed for a particular state shape; it does not claim that the underlying proof strategy is impossible. This distinction is important: negative memory guides exploration, but it does not rule out proof strategies as a matter of logic.

5. Coordinating Proof Decomposition and Lemma Proving

Sections 3 and 4 prove lemmas whose statements are fixed. We now start from only the specification and the final theorem, whose proof is left open, and let ShannonProver construct the decomposition as well (Phase II in Fig. 1). Phase II need not finish before Phase III starts: whether a decomposition is good is often known only after trying to prove its lemmas. An invariant may turn out to be too weak to carry through an oracle call, or a planned game hop may violate the restrictions on which modules the adversary may access. As experts do, ShannonProver moves back and forth between the two phases (Fig. 2).

Decomposer. The decomposer is an LLM agent that works on an EasyCrypt file containing the specification and the statement of the final theorem. Its goal is the same file with a complete proof: EasyCrypt must accept the whole file, with no `admit` remaining. In our experiments, the final theorem and the supplied definitions must remain unchanged, and new declarations may be added only in a marked region of the file, so the decomposer cannot make the theorem easier by weakening it or adding assumptions. The decomposer chooses the intermediate games, writes the auxiliary lemmas together with the modules and invariants they need, and assembles the proof of the final theorem; it may also prove lemmas directly. It checks its work by replaying the whole file in EasyCrypt, temporarily allowing `admit` placeholders, so that it can confirm that a skeleton of the argument type-checks and connects to the final theorem before every lemma is proved.

Handing lemmas to the lemma prover. When a lemma would benefit from step-by-step interaction with the proof state, the decomposer hands it to the lemma prover as a *job*. A job works on exactly one named lemma, using the proof-state compiler and the proof tree of §3 and §4. The decomposer never hands over the final theorem or asks the lemma prover to invent the overall reduction: choosing game hops and lemma statements remains its own responsibility. Each job receives a snapshot of the file, the name of the lemma, a time limit, and usually a short note with the goal, any proof prefix already accepted by EasyCrypt, the intended strategy, and the attempts that have already failed. A job that receives an accepted prefix first replays it, so it never starts from a state that the checker has not accepted. Based on checked progress rather than elapsed time, the decomposer requests one of three kinds of job: an exploratory job that tests whether a game hop, an invariant, or a proof strategy works; a completion job once accepted tactics already fix the strategy; or a continuation of an earlier job from its saved state. Handing off a lemma does not block the decomposer: a limited number of jobs run at the same time (two in our experiments), while the decomposer continues with independent parts of the argument, such as later games or probability bounds.

Feedback from Phase III to Phase II. A job that proves its lemma returns the proof, which is checked again against the

current file before it is merged. Otherwise, the job returns the longest prefix of tactics that EasyCrypt accepted, the remaining goals, and, when available, a saved state for a continuation; any notes it adds are hints, not checked facts. From such a result, the decomposer can finish the proof itself, request a continuation, restate or split the lemma or add helper lemmas, or change the decomposition. The last two are where Phase III feeds back into Phase II (Fig. 2; §6.4 describes an example from the ChaChaPoly case study).

6. Evaluation

Implementation. ShannonProver is an open-source research prototype written in Python and built on EasyCrypt. The core implementation contains about 89K lines of code: about 86K lines of Python and about 3K lines of OCaml that extend EasyCrypt. By component, it comprises about 33K lines for the multi-agent tree search (§4), including the orchestrator, the agents, and recovery from agent crashes; about 26K lines for the proof-state compiler (§3); about 24K lines for running EasyCrypt and parsing its output; and about 7K lines for the decomposer (§5). App. B describes the architecture in detail. Our code is available at https://github.com/SkyShannonProver/shannon-prover.

6.1. Evaluation Design

Our evaluation has two parts. The first is a controlled ablation of the lemma prover's two design ideas. For the proof-state compiler, we measure how much it reduces the agent's errors and the cost of recovering from them. For the multi-agent tree search, we measure how much it speeds up proof writing compared with a single agent. The second part is a case study on ChaCha20-Poly1305 [21], in which we evaluate both the lemma prover given the experts' decomposition and the end-to-end setting with the decomposer. Appendix C presents additional case studies on MEE-CBC [19] and CMAC [20], together with expert comments on these developments (App. C.3).

Setup. All ablations share the following setup. We use Opus 4.8[2] with high reasoning effort. The agent may read all project files and EasyCrypt source files, and we give it no hints about proof strategy. Each ablation compares the lemma prover with a baseline in which the mechanism under study is disabled (§6.2 and §6.3 describe each baseline). The baseline uses the same MCP-based EasyCrypt interaction layer as the lemma prover, so that differences in results do not stem from unrelated engineering factors such as session management and crash recovery.

We use the same stopping rule in all configurations: the agent may announce that it gives up at most three times. After each announcement, the system prompts the agent to try another strategy. If the agent makes three consecutive such announcements without proof progress, we count the run as a give-up. This rule prevents runs from spending an unbounded budget in states where the agent repeatedly reports that it is stuck. In pilot runs, forcing the agent to continue led either to wasted tokens at the same stuck point or to behavior unrelated to the proof, such as reading files outside the project.

Data sampling. To keep the API cost of the ablation study manageable while keeping the evaluation representative, we sample 40 lemmas from our benchmark dataset, covering different types and difficulty levels. Appendix C describes how we built this dataset. We group lemmas by proof shape: procedural (**P**), invariant-synthesis (**I**), and game-hop/reduction (**G**). This grouping is not meant to label cryptographic importance; rather, it separates qualitatively different proof-engineering regimes. Procedural lemmas are mostly *local*: once the relevant procedure body is exposed, progress is often made by a sequence of greedily discoverable symbolic-execution and simplification tactics. Invariant-synthesis lemmas require the agent to choose a relational invariant that is strong enough to carry through a call or loop boundary, but not so strong that it creates unnecessary side obligations. Game-hop and reduction lemmas require the proof to preserve higher-level game structure long enough to apply bridge lemmas, oracle equivalences, or probability bounds. These three shapes pose different difficulties for the agent, so aggregating them into a single average would obscure where the lemma prover helps.

Metrics. LLM-based theorem provers are usually evaluated by pass@k [26,27]: the model writes k complete proof scripts, and a theorem counts as solved if any of them is accepted. We do not use this metric, for two reasons. First, it suits short proofs that can be written in one shot, as in miniF2F [27]. Our lemmas often need dozens of tactics, so writing a complete proof in one shot that the checker accepts is unrealistic, and sampling many attempts (i.e., a very large k) is prohibitively expensive. Second, the binary outcome, whether a lemma is proved or not, does not show what our mechanisms change. For example, two runs may both succeed, one through a short, direct proof and the other only after many rejected tactics. We instead use metrics that target each mechanism, detailed in §6.2 and §6.3.

6.2. The Effects of the Proof-State Compiler

Baseline. The baseline has the agent interact directly with the checker: after submitting each tactic to the checker it sees only the resulting goal or the error message. The lemma prover instead shows the agent the output of the proof-state compiler (§3).

Metrics. A cheaper proof could come from two very different places, and a fair ablation has to separate them. The compiler's output might make the model propose fewer wrong tactics in the first place, or it might leave the first-attempt error rate unchanged while making those errors far less costly. We therefore have the following metrics:
❶ *Error-generation rate.* The fraction of the model's distinct proposed tactics that the checker rejects on its first attempt.

2. This was the strongest model available when we ran the ablation experiments.

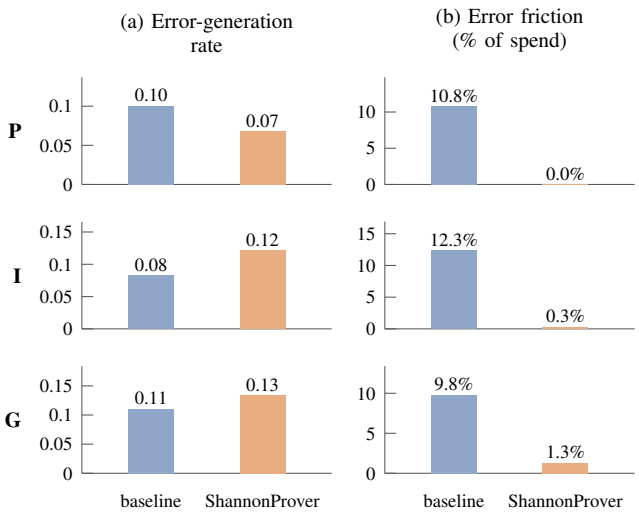

Figure 7. Compiler ablation on lemma pairs that both cases proved, grouped by proof shape: **P** = procedural, **I** = invariant-synthesis, **G** = game-hop/reduction.

❷ *Error friction.* The fraction of an agent's total API spend on sequences of consecutive tactics rejected by the checker. This captures the fraction that the agent spends on debugging syntax and recovering from errors.

To collect these metrics, our measurement pipeline separates deterministic metrics from qualitative judgement. The first stage is a Python parser: for each run it extracts the submitted tactics, the checker responses, and the agent reasoning traces for each submitted tactic. The second stage uses an LLM for cases a parser cannot classify mechanically: for example, when the agent commits several steps and later reverts them with no explicit checker error, so the parser cannot tell an unproductive dead end from a deliberate restructuring; there the LLM labels those tactic moves with the corresponding qualitative cause. The final stage is deterministic again: a script aggregates the parsed events and the labels into the reported numbers.

Results. Figure 7 shows that the proof-state compiler does not primarily work by making the model intrinsically less error-prone. The error-generation rate is close between the baseline and the lemma prover across the three lemma types. This is expected, because this metric is dominated by the model's own stochastic search behavior.

The error-friction numbers show the more important effect. Across all lemma types, the lemma prover sharply reduces the fraction of spend on such rejected sequences. This means that even when the agent still makes mistakes when it first tries a move, the agent can immediately recover from it. Meanwhile, the agent in the baseline must figure out the right move via trial-and-error. The ablation shows that the proof-state compiler is not a proof oracle that provides the right move, but an interface that provides sufficient information to free the agent from mechanical debugging.

The main counter-intuitive result for the error-generation rate is invariant-synthesis lemmas, where the lemma prover has a visibly higher rate. The experimental traces suggest why: with the richer information provided by the compiler, the agent often attempts stronger and more structured invariants, because it can see more of the resources that apply, the program statements to which they apply, and the required bindings. Some of these ambitious invariant attempts are rejected on the first attempt. In the baseline, however, the agent is more likely to choose a weaker invariant that passes immediately, even if it leaves the proof in a worse state later.

The compiler ablation therefore shows that the proof-state compiler mainly reduces the cost of errors rather than their number: the agent still makes mistakes, but it spends far less of its budget recovering from them.

6.3. The Effects of Multi-Agent Tree Search

Baseline. The multi-agent tree search has two layers (§4): backtracking within a branch, and spawning further branches from shared checkpoints. Backtracking itself cannot be meaningfully removed: an agent that cannot undo a wrong step can only fail. The baseline therefore keeps backtracking and removes spawning. It runs one branch at a time: the agent can undo and backtrack within its branch, and if its agent process crashes, the branch is resumed from its committed prefix, but no sibling branch ever races it. The lemma prover adds the cross-agent tree policy: a second branch can race from a shared checkpoint, and dead or stalled nodes can be respawned with negative memory so that the next agent does not repeat the same failed strategy. In both configurations, the agent can use the proof-state compiler.

Metrics. We measure how often the agent needs to backtrack: the fraction of solved runs that backtrack at least once, and the number of backtrack events per solved run. These numbers show how far proof search departs from a straight line in each lemma category. We also measure the solve rate in each category, and how it grows over the course of a run.

Results. Figure 8a–b first shows that the need for backtracking is highly lemma-type dependent. Procedural lemmas are almost straight-line, as expected from their local structure. In contrast, invariant-synthesis and game-hop lemmas are much less linear. 38% of solved invariant runs and 43% of solved game-hop runs backtrack at least once, with about 2.0–2.3 backtrack events per solved run. These are exactly the lemmas where an attempted tactic that is accepted by the checker can lead the agent into an unproductive strategy: an invariant may be too weak or too strong, and a game-hop proof may descend too early and lose a bridge lemma or oracle-equivalence resource. For these lemmas, the ability to rewind is not a rare recovery feature; it is part of how successful proofs are found.

Figure 8c–d shows what spawning adds on top of an agent that backtracks. On procedural lemmas, backtracking within one branch already suffices: the single agent proves every lemma, and a second branch only wins the race. On invariant-synthesis and game-hop lemmas, it does not: even with backtracking, a third of single-agent runs fail,

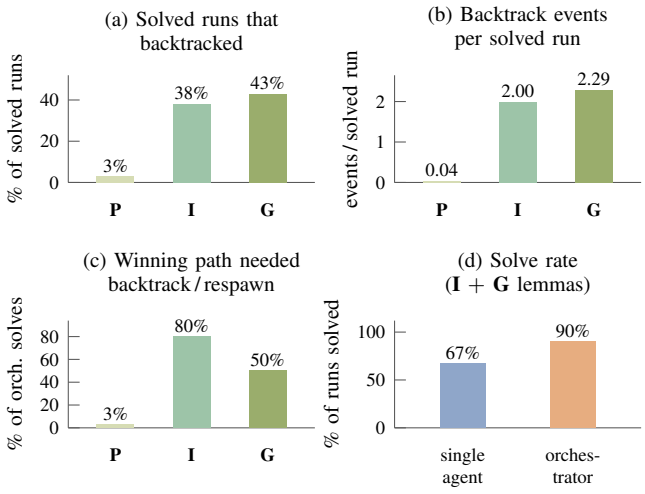

Figure 8. Tree-search ablation, grouped by proof shape (**P** = procedural; **I** = invariant-synthesis; **G** = game-hop/reduction).

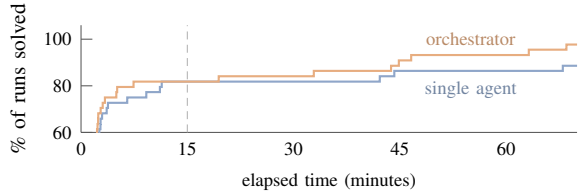

Figure 9. Fraction of lemmas solved by elapsed time. The two configurations coincide for the first 15 minutes (dashed line) and separate afterwards.

and spawning raises the solve rate from 67% to 90%. The gain comes mainly from redundancy: the second branch completed the proof after the root branch's agent had died. It also appears late (Fig. 9): both configurations solve 82% of all runs within 15 minutes and separate only afterwards, once a first branch has stalled.

The tree ablation therefore supports two conclusions. First, proof search is not uniformly non-linear: for local procedural obligations, a single branch with backtracking is enough. Second, for the proof types that matter most in cryptographic developments—invariants and game hops—backtracking within one branch is not enough, and spawning further branches makes the search robust to branches that stall or die.

6.4. Case Studies

ChaCha20-Poly1305 [21] is a representative case study project. It is an AEAD construction that combines the ChaCha20 stream cipher with the Poly1305 authenticator and is used in IETF protocols and deployed in TLS [5]. Its EasyCrypt formal proof contains a mixture of low-level algebraic reasoning, invariant-heavy relational proofs, and game-hop/reduction lemmas. Experts described the development as interleaving the design of new abstractions with mechanical proof engineering: abstraction choices shaped the proof structure, while mechanization difficulties fed back into the abstraction design.

This case study demonstrates that LLM agents today can fully automate lemma-level proof (Phase III in Fig. 1) at the scale and complexity of ChaChaPoly. Figure 10 shows all the lemmas in ChaChaPoly and the API cost per lemma. The procedural lemmas are generally easy and they are solved at low API cost. This is already useful because large projects contain many such lemmas, and writing them manually still consumes expert time. Importantly, ShannonProver solved the hardest lemmas in the project, while the baseline agent failed. These hard lemmas are fewer in the project, but they are also the most important: they are the points where the lemma decomposition is tested. If these lemmas can be proven, the expert gains evidence that the decomposition has the right structure. If LLM agents or humans fail to prove a lemma, the thought process in the failed attempts is still informative: it often indicates that the current lemma is missing an invariant, that a game hop has not been modularized at the right boundary, or that some other part of the decomposition needs to be revised. This feedback is valuable for redoing the decomposition.

We also evaluate proof quality on three fronts: concision (**C**), maintainability/readability (**M**) and semantic proof idiom (**S**). **C** measures whether the proof is compact and avoids unnecessary tactic churn or duplicated reasoning. **M** measures whether the proof structure is easy for a human to inspect, debug, and adapt when the surrounding definitions change. **S** measures whether the proof follows the natural EasyCrypt proof structure and uses existing proven facts, rather than reproving low-level details. We have two experts independently receive each pair of proofs in randomized order and manually evaluate for each dimension. Figure 10 shows that the AI-generated proofs still leave room for improvement on invariant and game-hop lemmas, making proof quality a promising direction for future work.

The LLM API cost is also reasonable to make iterative proof development practical. Across the ChaCha20-Poly1305 lemmas (Fig. 10), proving all lemmas of the expert decomposition once costs $233.65 for API calls and completes within a day of wall-clock time. The cost is concentrated in the harder lemmas: the 27 procedural lemmas together cost $40.98, while the invariant-synthesis and game-hop/reduction lemmas account for $76.99 and $115.68, respectively. This distribution is desirable from a workflow perspective: the system spends little on procedural lemmas, and most of the cost goes to the lemmas that are most informative for testing the decomposition. Even if a cryptographer revises the decomposition and reruns the proof pass several times, the marginal cost remains on the order of a few hundred dollars per iteration and the time is at most a day, rather than person-weeks or person-months of expert proof engineering. This makes ShannonProver useful not only as a one-shot proof generator, but also as a practical tool for iterating on large cryptographic formalizations.

Evaluation for end-to-end proof. We further evaluate ShannonProver in an end-to-end setting, where the system is given only the specification and the statement of the main theorem. In this workflow, the complete proof was

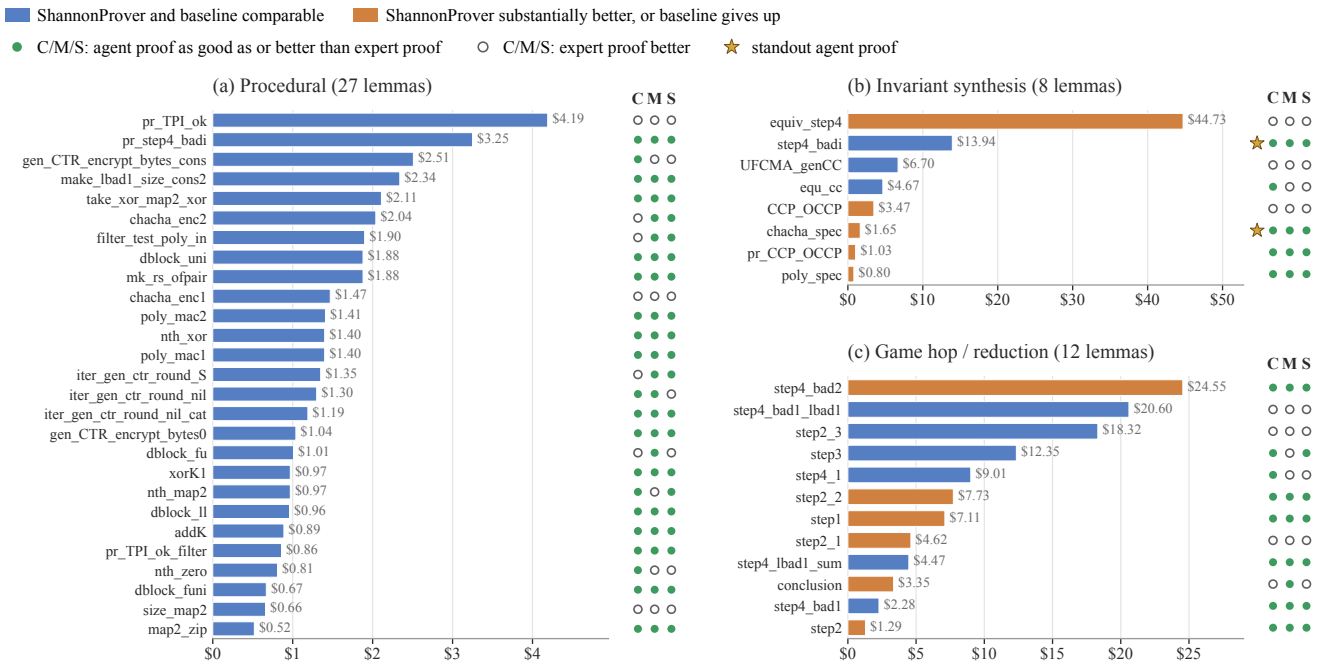

Figure 10. Per-lemma API cost for the ChaCha20-Poly1305 project, grouped by proof shape (**P/I/G**). Bar color compares ShannonProver with the baseline. For concision (**C**), maintainability/readability (**M**), and semantic proof idiom (**S**), a filled green dot denotes an agent proof comparable to or better than the expert proof, while a hollow gray circle denotes an expert advantage. Stars mark the two agent proofs described in Appendix D.

constructed in 5 hours, using a strong model (Opus 5[3]) as decomposer and keep the lemma provers same as before. In this run, about 20% of the lemmas in the decomposition were reworked after the decomposer first wrote them and received feedback from the lemma prover. The total cost is about $420 for a complete proof. In comparison, a baseline that uses CLI with Opus 5 takes 8 hours to complete the proof end to end; and Opus 4.8 failed to complete the task.

7. Related Work

Machine-checked cryptography. Tools for machine-checked cryptography trade expressiveness for automation. Symbolic tools such as Tamarin [28] and ProVerif [29] are the most automatic: they model the adversary as a Dolev-Yao attacker [30] and treat cryptographic primitives as perfect, so they cannot state how likely an attack succeeds. CryptoVerif [31] brings automation to the computational model, but it builds proofs from a fixed set of game transformations and library assumptions; arguments outside this set, such as a new reduction or reasoning about the mathematics of a scheme, need extensive manual guidance or cannot be carried out. Proof assistants such as EasyCrypt [18] and others [32–36] can express any reduction together with the mathematics it relies on, and experts have used them to verify standardized schemes and NIST post-quantum standards [14,19,21,22,37,38], in some cases down to verified implementations [9,39,40]; in exchange, the user writes every proof step. LLM agents change this trade-off: they can carry out the many detailed steps that a proof assistant leaves to the user, so automation no longer requires giving up expressiveness.

Among proof assistants, we build on EasyCrypt: it supports the full range of reasoning in game-based proofs, and the expert developments in our case studies were originally written in it. Lean has a much larger ecosystem and is the target of most LLM-based theorem proving, but game-based proofs there first need a library for probabilistic programs with oracles. VCVio [36] provides such a library, with oracle computations and a relational program logic; it is a promising foundation, but it is still under active development, and its core definitions and semantics change frequently. We plan to port ShannonProver to Lean as such libraries mature.

LLMs for formal proofs. LLM-based theorem proving has advanced rapidly, mostly in mathematics, through whole-proof generation [26,41–45] evaluated on competition benchmarks [27,46,47], tactic-level interaction with the proof assistant [23,48,49], and, more recently, agentic provers [50,51]. In cryptography, reinforcement learning and learned search heuristics have guided symbolic protocol proofs [52,53], LLMs have been used to build or check protocol models [54,55], generative AI has assisted a researcher in formalizing computational security proofs in Lean [56], and AI provers have discharged correctness obligations extracted from cryptographic Rust code [57]. Other work enriches the interface between model and prover, by

3. We ran this experiment during the revision of this paper, after the other ablations; Opus 5 was the strongest model available at that time.

suggesting tactics in the human workflow [58,59], exposing proof states to agents [60], or adding a natural-language representation of the proof state [61].

8. Conclusion and Future Work

This work shows that LLM agents can take on a substantial part of mechanized cryptographic proof when the system around them supplies the right context and controls the search.

Two problems remain open. Decomposing a large proof obligation into provable lemmas (Phase II) still depends strongly on the capability of the model; better harness support for the decomposer may improve performance. For Phase I, it is interesting to explore how AI can help experts design the specification itself, or help the experts design the abstractions needed for certain cryptographic schemes.

9. Generative AI Responsibility Policy

Content and data holder rights. ShannonProver is evaluated on a benchmark of EasyCrypt developments comprising ML-KEM, MEE-CBC, ChaCha20-Poly1305, and CMAC. We used each artifact in a manner consistent with the terms under which its rights holders released it.

The ML-KEM specification is released under the Creative Commons CC0 1.0 Universal public domain dedication. The MEE-CBC and ChaCha20-Poly1305 developments are distributed with the EasyCrypt proof assistant, under the MIT License.

The CMAC development is not publicly available. We obtained explicit authorization from its authors to evaluate ShannonProver on it, and our use is confined to the scope of that authorization. We therefore report aggregate results for CMAC without releasing the underlying source files; readers seeking access should contact the original authors directly.

Environmental footprint. ShannonProver substantially reduces token usage compared with using Codex or Claude through a plain CLI to solve the same task. We believe this points to a more responsible way of using LLMs for proof automation going forward: whenever possible, work that does not require substantial model reasoning should be moved to deterministic components, reserving LLM calls for tasks where their reasoning capabilities are genuinely needed.

Acknowledgments

This work is supported by gifts from Accenture, Algorithmic SuperIntelligence Labs, Amazon, AMD, Anyscale, Broadcom, cmpnd, Google, IBM, Intel, Intesa Sanpaolo, Lambda, Lightspeed, Mirendil, NVIDIA, Samsung SDS, and VESSL.

We would like to thank Quang Dao, Shafi Goldwasser, Mohsen Lesani, Hongyi Ling, Ning Luo, Surya Mathialagan, Yibin Yang, the students in the security group at Sky Lab, and the participants and organizers in HACS 2026 for valuable discussions. We also thank Sebastian Angel and AWS Cryptography Group, especially Tal Rabin, Yuval Ishai, and Shai Halevi, for valuable discussion and feedback on this draft and future directions.

Competing Interests

Apart from the research funding listed in the acknowledgments, the authors declare no competing financial or non-financial interests.

References

[1] S. Halevi, "A plausible approach to computer-aided cryptographic proofs," Cryptology ePrint Archive, Paper 2005/181, 2005.

[2] National Institute of Standards and Technology, "FIPS 202: SHA-3 standard: Permutation-based hash and extendable-output functions," Federal Information Processing Standards Publication 202, 2015.

[3] National Institute of Standards and Technology, "FIPS 197: Advanced encryption standard (AES)," Federal Information Processing Standards Publication 197, 2023.

[4] National Institute of Standards and Technology, "FIPS 203: Module-lattice-based key-encapsulation mechanism standard," Federal Information Processing Standards Publication 203, 2024.

[5] E. Rescorla, "The transport layer security (TLS) protocol version 1.3," RFC 8446, 2018.

[6] D. J. Bernstein, T. Lange, and P. Schwabe, "The security impact of a new cryptographic library," in *Progress in Cryptology – LATINCRYPT 2012*, ser. Lecture Notes in Computer Science, vol. 7533. Springer, 2012, pp. 159–176.

[7] libsodium contributors, "libsodium documentation," https://doc.libsodium.org/, accessed 2026-05-09.

[8] OpenSSL Software Foundation, "OpenSSL: Cryptography and SSL/TLS toolkit," https://www.openssl.org/, accessed 2026-05-09.

[9] J. Protzenko, B. Parno, A. Fromherz, C. Hawblitzel, M. Polubelova, K. Bhargavan, B. Beurdouche, J. Choi, A. Delignat-Lavaud, C. Fournet, N. Kulatova, T. Ramananandro, A. Rastogi, N. Swamy, C. M. Wintersteiger, and S. Zanella-Béguelin, "EverCrypt: A fast, verified, cross-platform cryptographic provider," in *2020 IEEE Symposium on Security and Privacy*. IEEE, 2020, pp. 983–1002.

[10] M. Thomson and S. Turner, "Using TLS to secure QUIC," RFC 9001, 2021.

[11] K. Cohn-Gordon, C. J. F. Cremers, B. Dowling, L. Garratt, and D. Stebila, "A formal security analysis of the Signal messaging protocol," in *2017 IEEE European Symposium on Security and Privacy (EuroS&P)*. IEEE, 2017, pp. 451–466.

[12] J. A. Donenfeld, "WireGuard: Next generation kernel network tunnel," in *24th Annual Network and Distributed System Security Symposium (NDSS 2017)*. Internet Society, 2017.

[13] J. B. Almeida, S. A. Olmos, M. Barbosa, G. Barthe, F. Dupressoir, B. Grégoire, V. Laporte, J.-C. Léchenet, C. Low, T. Oliveira, H. Pacheco, M. Quaresma, P. Schwabe, and P.-Y. Strub, "Artifact for formally verifying kyber episode V: Machine-checked IND-CCA security and correctness of ML-KEM in EasyCrypt," IACR Artifact Archive, crypto/2024/a3, 2024.

[14] M. Barbosa, F. Dupressoir, A. Hülsing, M. Meijers, and P.-Y. Strub, "A tight security proof for SPHINCS+, formally verified," in *Advances in Cryptology – ASIACRYPT 2024*, ser. Lecture Notes in Computer Science, vol. 15487. Springer, 2024, pp. 35–67.

[15] Y. Bertot and P. Castéran, *Interactive Theorem Proving and Program Development: Coq'Art: The Calculus of Inductive Constructions*, ser. Texts in Theoretical Computer Science. Springer, 2004.

[16] L. de Moura and S. Ullrich, "The Lean 4 theorem prover and programming language," in *Automated Deduction – CADE 28*, ser. Lecture Notes in Computer Science, vol. 12699. Springer, 2021, pp. 625–635.

[17] N. Swamy, C. Hriţcu, C. Keller, A. Rastogi, A. Delignat-Lavaud, S. Forest, K. Bhargavan, C. Fournet, P.-Y. Strub, M. Kohlweiss, J.-K. Zinzindohoué, and S. Zanella-Béguelin, "Dependent types and multi-monadic effects in F*," in *Proceedings of the 43rd Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL)*. ACM, 2016, pp. 256–270.

[18] G. Barthe, B. Grégoire, S. Heraud, and S. Z. Béguelin, "Computer-aided security proofs for the working cryptographer," in *Advances in Cryptology – CRYPTO 2011*, ser. Lecture Notes in Computer Science, vol. 6841. Springer, 2011, pp. 71–90.

[19] J. B. Almeida, M. Barbosa, G. Barthe, and F. Dupressoir, "Verifiable side-channel security of cryptographic implementations: Constant-time MEE-CBC," in *Fast Software Encryption (FSE)*, 2016.

[20] C. Baritel-Ruet, F. Dupressoir, P.-A. Fouque, and B. Grégoire, "Formal security proof of cmac and its variants," in *2018 IEEE 31st Computer Security Foundations Symposium (CSF)*. IEEE, 2018, pp. 91–104.

[21] J. B. Almeida, M. Barbosa, G. Barthe, B. Grégoire, A. Koutsos, V. Laporte, T. Oliveira, and P.-Y. Strub, "The last mile: High-assurance and high-speed cryptographic implementations," in *IEEE Symposium on Security and Privacy*, 2020.

[22] J. B. Almeida, S. A. Olmos, M. Barbosa, G. Barthe, F. Dupressoir, B. Grégoire, V. Laporte, J.-C. Léchenet, C. Low, T. Oliveira, H. Pacheco, M. Quaresma, P. Schwabe, and P.-Y. Strub, "Formally verifying kyber episode V: Machine-checked IND-CCA security and correctness of ML-KEM in EasyCrypt," in *Advances in Cryptology – CRYPTO 2024*, ser. Lecture Notes in Computer Science, vol. 14921. Springer, 2024, pp. 384–421.

[23] K. Yang, A. Swope, A. Gu, R. Chalamala, P. Song, S. Yu, S. Godil, R. J. Prenger, and A. Anandkumar, "Leandojo: Theorem proving with retrieval-augmented language models," *Advances in Neural Information Processing Systems*, vol. 36, pp. 21 573–21 612, 2023.

[24] S. Rajaee, K. Pratik, G. Cesa, and A. Behboodi, "Local look-ahead guidance via verifier-in-the-loop for automated theorem proving," in *Findings of the Association for Computational Linguistics: ACL 2025*, 2025.

[25] A. V. Aho, M. S. Lam, R. Sethi, and J. D. Ullman, *Compilers: Principles, Techniques, and Tools*, 2nd ed. Addison-Wesley, 2006.

[26] S. Polu and I. Sutskever, "Generative language modeling for automated theorem proving," *arXiv preprint arXiv:2009.03393*, 2020.

[27] K. Zheng, J. M. Han, and S. Polu, "MiniF2F: A cross-system benchmark for formal olympiad-level mathematics," in *ICLR*, 2022.

[28] S. Meier, B. Schmidt, C. Cremers, and D. Basin, "The TAMARIN prover for the symbolic analysis of security protocols," in *Computer Aided Verification*, ser. Lecture Notes in Computer Science, vol. 8044. Springer, 2013, pp. 696–701.

[29] B. Blanchet, "Modeling and verifying security protocols with the applied pi calculus and ProVerif," *Foundations and Trends in Privacy and Security*, vol. 1, no. 1–2, pp. 1–135, 2016.

[30] D. Dolev and A. C. Yao, "On the security of public key protocols," *IEEE Transactions on Information Theory*, vol. 29, no. 2, pp. 198–208, 1983.

[31] B. Blanchet, "A computationally sound mechanized prover for security protocols," *IEEE Transactions on Dependable and Secure Computing*, 2008.

[32] G. Barthe, B. Grégoire, and S. Z. Béguelin, "Formal certification of code-based cryptographic proofs," in *POPL*, 2009.

[33] D. A. Basin, A. Lochbihler, and S. R. Sefidgar, "CryptHOL: Game-based proofs in higher-order logic," *Journal of Cryptology*, 2020.

[34] P. G. Haselwarter, E. Rivas, A. Van Muylder, T. Winterhalter, C. Abate, N. Sidorenco, C. Hriţcu, K. Maillard, and B. Spitters, "Ssprove: A foundational framework for modular cryptographic proofs in coq," *ACM transactions on programming languages and systems*, vol. 45, no. 3, pp. 1–61, 2023.

[35] A. Petcher and G. Morrisett, "The foundational cryptography framework," in *Principles of Security and Trust (POST)*, 2015.

[36] D. Tuma, Q. Dao, J. Waters, A. Hicks, and N. Hopper, "Vcvio: Verified cryptography in lean via oracle effects and handlers," *Cryptology ePrint Archive*, 2026.

[37] M. Barbosa, F. Dupressoir, B. Grégoire, A. Hülsing, M. Meijers, and P.-Y. Strub, "Machine-checked security for XMSS as in RFC 8391 and SPHINCS+," in *Advances in Cryptology – CRYPTO 2023*, ser. Lecture Notes in Computer Science, vol. 14085. Springer, 2023, pp. 421–454.

[38] J. B. Almeida, C. Baritel-Ruet, M. Barbosa, G. Barthe, F. Dupressoir, B. Grégoire, V. Laporte, T. Oliveira, A. Stoughton, and P.-Y. Strub, "Machine-checked proofs for cryptographic standards: Indifferentiability of sponge and secure high-assurance implementations of SHA-3," in *ACM CCS*, 2019.

[39] J.-K. Zinzindohoué, K. Bhargavan, J. Protzenko, and B. Beurdouche, "HACL*: A verified modern cryptographic library," in *ACM CCS*, 2017.

[40] J. B. Almeida, M. Barbosa, G. Barthe, A. Blot, B. Grégoire, V. Laporte, T. Oliveira, H. Pacheco, B. Schmidt, and P.-Y. Strub, "Jasmin: High-assurance and high-speed cryptography," in *ACM CCS*, 2017.

[41] G. Lample, T. Lacroix, M.-A. Lachaux, A. Rodriguez, A. Hayat, T. Lavril, G. Ebner, and X. Martinet, "Hypertree proof search for neural theorem proving," *Advances in neural information processing systems*, vol. 35, pp. 26 337–26 349, 2022.

[42] H. Xin *et al.*, "DeepSeek-Prover-V1.5: Harnessing proof assistant feedback for reinforcement learning and monte-carlo tree search," *arXiv preprint arXiv:2408.08152*, 2024.

[43] DeepSeek-AI, "DeepSeek-Prover-V2: Advancing formal mathematical reasoning via reinforcement learning for subgoal decomposition," *arXiv preprint arXiv:2504.21801*, 2025.

[44] Y. Lin, S. Tang, B. Lyu, J. Wu, H. Lin, K. Yang, J. LI, M. Xia, D. Chen, S. Arora *et al.*, "Goedel-prover: A frontier model for open-source automated theorem proving," in *Second Conference on Language Modeling*.

[45] T. Hubert, R. Mehta, L. Sartran, M. Z. Horváth, G. Žužić, E. Wieser, A. Huang, J. Schrittwieser, Y. Schroecker, H. Masoom *et al.*, "Olympiad-level formal mathematical reasoning with reinforcement learning," *Nature*, pp. 1–3, 2025.

[46] Z. Azerbayev, B. Piotrowski, H. Schoelkopf, E. W. Ayers, D. Radev, and J. Avigad, "ProofNet: Autoformalizing and formally proving undergraduate-level mathematics," *arXiv preprint arXiv:2302.12433*, 2023.

[47] G. Tsoukalas, J. Lee, J. Jennings, J. Xin, M. Ding, M. Jennings, A. Thakur, and S. Chaudhuri, "Putnambench: Evaluating neural theorem-provers on the putnam mathematical competition," *Advances in Neural Information Processing Systems*, vol. 37, pp. 11 545–11 569, 2024.

[48] A. Q. Jiang, W. Li, S. Tworkowski, K. Czechowski, T. Odrzygóźdź, P. Miłoś, Y. Wu, and M. Jamnik, "Thor: Wielding hammers to integrate language models and automated theorem provers," *Advances in Neural Information Processing Systems*, vol. 35, pp. 8360–8373, 2022.

[49] E. First, M. N. Rabe, T. Ringer, and Y. Brun, "Baldur: Whole-proof generation and repair with large language models," in *Proceedings of the 31st ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering*, 2023, pp. 1229–1241.

[50] B. Breen, M. D. Tredici, J. McCarran *et al.*, "Ax-Prover: A deep reasoning agentic framework for theorem proving in mathematics and quantum physics," *arXiv preprint arXiv:2510.12787*, 2025.

[51] A. Sosso, A. Arora, and B. Spitters, "Agentic proving for program verification," *arXiv preprint arXiv:2605.23772*, 2026.

[52] Y. Xiong, C. Su, W. Huang, F. Miao, W. Wang, and H. Ouyang, "SmartVerif: Push the limit of automation capability of verifying security protocols by dynamic strategies," in *29th USENIX Security Symposium (USENIX Security 20)*, 2020, pp. 253–270.

[53] M. Cosler, C. Cremers, B. Finkbeiner, M. Ghanem, and N. Medinger, "Less effort, shorter proofs: Reinforcement learning for security protocol analysis in Tamarin," *arXiv preprint arXiv:2605.23643*, 2026.

[54] C. Curaba, D. D'Ambrosi, A. Minisini, and N. P.-C. Antolín, "CryptoFormalEval: Integrating LLMs and formal verification for automated cryptographic protocol vulnerability detection," *arXiv preprint arXiv:2411.13627*, 2024.

[55] Z. Mao, J. Wang, J. Sun, S. Qin, and J. Xiong, "LLM-aided automatic modeling for security protocol verification," in *International Conference on Software Engineering (ICSE)*, 2025.

[56] B. Spitters, "CatCrypt: From Rust to cryptographic security in Lean," *Cryptology ePrint Archive, Paper 2026/604*, 2026.

[57] N. Klaus, J. Conejero, and P. Tolmach, "A Rust-to-Lean verification pipeline with AI provers: An experience report," *arXiv preprint arXiv:2605.30106*, 2026.

[58] P. Song, K. Yang, and A. Anandkumar, "Lean Copilot: Large language models as copilots for theorem proving in Lean," *arXiv preprint arXiv:2404.12534*, 2024.

[59] S. Welleck and R. Saha, "LLMSTEP: LLM proofstep suggestions in Lean," *arXiv preprint arXiv:2310.18457*, 2023.

[60] L. Aniva, C. Sun, B. Miranda, C. Barrett, and S. Koyejo, "Pantograph: A machine-to-machine interaction interface for advanced theorem proving, high level reasoning, and data extraction in Lean 4," *arXiv preprint arXiv:2410.16429*, 2024.

[61] L. Teodorescu, G. Baudart, E. J. G. Arias, and M. Lelarge, "NLIR: Natural language intermediate representation for mechanized theorem proving," *NeurIPS MATH-AI Workshop*, 2024.

[62] J. B. Almeida, M. Barbosa, G. Barthe, B. Grégoire, V. Laporte, J.-C. Léchenet, T. Oliveira, H. Pacheco, M. Quaresma, P. Schwabe *et al.*, "Formally verifying kyber: Episode iv: Implementation correctness," *IACR Transactions on Cryptographic Hardware and Embedded Systems*, vol. 2023, no. 3, pp. 164–193, 2023.

Appendix A. Deferred Explanations of Concepts

A.1. Proof Resources

A *proof resource* is something the proof can use instead of reasoning from scratch. The main kinds are lemmas already proved in the project or in a library (including equivalence lemmas, bridge lemmas, and probability bounds), axioms and assumptions of the specification, and hypotheses in the current proof. Each of them packages reasoning that has already been done or assumed; using it at the right point saves the proof from redoing that reasoning.

Two properties of resources matter for ShannonProver. First, resources are chosen, not invented: the agent decides which resource to use, whereas an invariant or a new lemma has to be created before it can be used. Second, which resources can be used changes as the proof proceeds. Choosing a resource is a strategic decision left to the agent; finding the resources that apply at the current state and instantiating them is mechanical work done by the compiler (§3). Filling a premise of the chosen resource with a hypothesis is part of this mechanical work: when a hypothesis in the current proof states exactly that premise, the compiler uses it (§3.4).

Table 1 places proof resources among the other objects that appear in a proof step, using the running example of §3.

TABLE 1. OBJECTS IN A PROOF STEP AND HOW SHANNONPROVER TREATS THEM. EXAMPLES REFER TO THE RUNNING EXAMPLE OF §3.

Object	Example	Role in a proof step	Handled by
Proof resource: lemma, axiom, local hypothesis	`Alossless_F`	Used instead of new reasoning	Agent chooses; compiler finds and instantiates
Hypothesis stating a premise of the chosen resource	`Ofll : islossless D2(O).O.f`	Fills that premise	Compiler
Invariant, intermediate game, new lemma	A relational invariant for a loop	Must be invented	Agent
Module, type	`D2(O).O` for the parameter `X`	Fills an argument of a resource	Compiler
Program structure	The call to `guess` in the goal	Decides whether a resource applies now	Compiler tracks it
Side condition	`islossless D2(O).O.f`, if no hypothesis states it	A goal that remains to be proved	Agent proves it

Program structure decides which resources apply. Program structure is not a resource, but it determines which resources can be used. In game-based proofs, many resources apply only at particular program points: a lemma about a procedure call applies only while that call is still visible in the goal, and a lemma about an oracle applies only at a call to that oracle. Each accepted tactic changes the visible program structure, so a resource may not apply yet, apply now, or no longer apply. For example, after `inline*` replaces a call with the body of the procedure, a lemma about that call no longer applies to the resulting goals; this is the unconscious lowering of §2. For this reason, the compiler recomputes at every proof state which resources apply, instead of treating the project's lemmas as a fixed library.

What the agent sees. When the compiler shows a resource to the agent, it gives the name of the resource, its statement instantiated at the current state, and which of its arguments remain open (§3.3).

Appendix B. System Architecture

Figure 11 shows how ShannonProver is organized. The architecture follows one rule: an LLM agent may choose proof actions, but it does not own the proof state and does not decide when a proof is valid. The figure has four bands, which also define the code categories reported in Section 6.

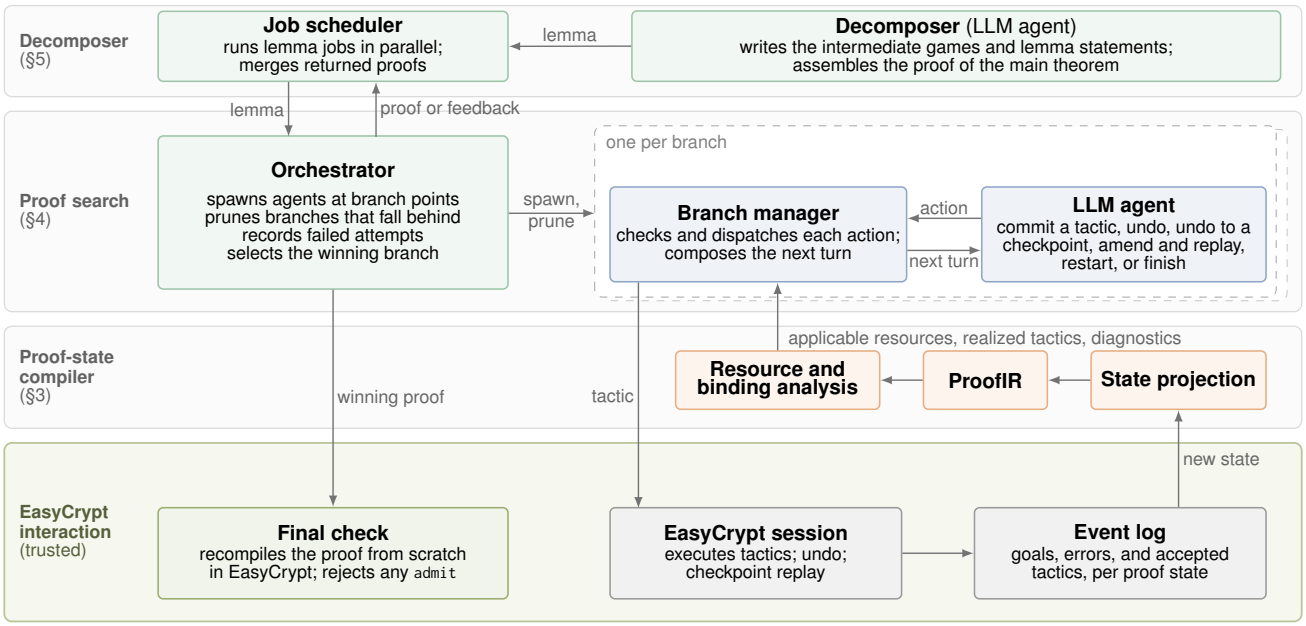

Figure 11. ShannonProver's architecture. The decomposer hands each lemma to the lemma prover as a job. Within a job, the orchestrator runs one LLM agent per branch of the proof tree: the agent chooses an action, the branch manager dispatches it to an EasyCrypt session, and the proof-state compiler turns the resulting state into the agent's next turn. Only the bottom band is trusted: a tactic changes the proof only if EasyCrypt accepts it, and a proof counts only after the final check recompiles it from scratch.

Decomposer. The decomposer is the LLM agent of Section 5. It writes the intermediate games and lemma statements and submits individual lemmas as jobs to a scheduler, which runs a bounded number of jobs in parallel. Each job runs the lemma prover on one lemma and returns either a proof that has passed the final check or feedback for the decomposer.

Proof search. Within a job, the orchestrator manages the proof tree of Section 4. Each branch is extended by one LLM agent, paired with a branch manager. The agent acts through six actions: commit a tactic, undo the last step, undo to a checkpoint, amend an earlier tactic and replay, restart, or finish. The branch manager checks each action, dispatches it to an EasyCrypt session, and composes the agent's next turn from the current goal, the outcome of the last action, and the output of the proof-state compiler. The orchestrator observes all branches: it spawns agents at branch points, prunes branches that fall behind, records failed attempts, and selects the winning branch.

Proof-state compiler. After each action that changes the proof state, the compiler reads the new state from the event log and runs the pipeline of Section 3: state projection, ProofIR, and resource and binding analysis. Its output, namely the applicable resources, realized tactics, and diagnostics, becomes part of the agent's next turn. The compiler only reads the proof state; it never changes it.

EasyCrypt interaction. The bottom band is the only part of the system that interacts with EasyCrypt, and it is the trust boundary. An EasyCrypt session executes tactics, supports undo and checkpoint replay, and records every goal, error, and accepted tactic in an event log. A tactic changes the proof only if EasyCrypt accepts it. When a branch finishes, the final check writes the proof into the source file, recompiles the file from scratch in EasyCrypt, and rejects the proof if any `admit` remains. Errors in the agents, the orchestrator, or the compiler can therefore waste time, but they cannot cause an invalid proof to be accepted.

Appendix C. Dataset and Additional Case Studies

C.1. Dataset Construction

Our benchmark collects proof tasks from EasyCrypt developments of cryptographic schemes. It complements mathematical theorem-proving resources such as LeanDojo [23] with tasks involving probabilistic and relational programs, adversaries and oracles, bad events, game hops, and security reductions. We use a subset spanning different proof types and difficulty levels for the ablation study in the main body. The public and private parts of the benchmark are described below.

Building the dataset. Each task in the benchmark is a single lemma taken from a real cryptographic development in EasyCrypt, together with the surrounding context it needs (i.e., sibling definitions, cloned theories, and imported lemmas). We remove the original proof; the task is to reconstruct a proof script for the lemma, with its statement and context unchanged. Each lemma is labeled with the scheme from which it comes, the cryptographic property it concerns, the proof techniques it requires, and a difficulty level.

The four parts and what each one measures. The benchmark is organized into four parts of increasing realism and difficulty. They differ in where the proofs come from, how hard they are, and crucially whether they could have been seen during a model's pretraining.

❶ *Foundational primitives.* Textbook and building-block schemes such as ElGamal and hashed-ElGamal, Schnorr, pseudorandom generators, private information retrieval, Pedersen commitments, and the correctness of Cramer–Shoup. These are short, local proofs, and they test whether the basic proof-state representation is adequate: finite-set and list algebra, simple relational coupling, losslessness, and single textbook reductions. This part establishes the floor of the benchmark and isolates interface issues from genuine reasoning difficulty.

❷ *Deployed real-world schemes.* This part includes standardized real-world schemes (e.g., ChaCha20-Poly1305, which we use as a case study). Unlike Part ❶, these are large, multi-lemma developments, so they test whether an LLM agent can move beyond isolated lemmas and find useful proof resources across lemmas, such as long game-hop chains, oracle invariants, bad-event probability bounds, and reasoning that depends on context established by other lemmas in the same project.

❸ *Post-quantum standards.* Lemmas from the machine-checked proofs of the NIST post-quantum standards: the IND-CCA security of ML-KEM (FIPS-203) [22,62], the tight security proof of SLH-DSA/SPHINCS$^+$ [14], the related XMSS development [37], and the security argument of ML-DSA/Dilithium. These are the most demanding and most current proofs in the benchmark, addressing the same problems that the standardization community itself verifies: lattice hardness assumptions, the Fujisaki-Okamoto transform, and hash-based signature arguments. Because these developments are recent and specialized, they also test whether an LLM agent can operate on proofs that lie well outside the well-trodden textbook material.

❹ *A published case study with non-public proof scripts.* This part consists of the CMAC/CBC-MAC family of proofs. We obtained the scripts directly from their authors, who authorized their use in this evaluation. To our knowledge, these scripts have not been publicly released. This reduces the chance that a model has seen these particular proofs during training, although it does not rule out exposure to related mathematics or proof techniques.

Sources and availability. Parts ❶–❸ are built from publicly available formal developments: the EasyCrypt example suite and standard libraries for the foundational and deployed schemes, and the public machine-checked artifacts for the post-quantum standards. For these parts, we will openly release our benchmark. Part ❹ is the exception: the CMAC development is private. Its proof scripts are not available in any public repository, and we obtained them directly from the development's authors for the purpose of this benchmark. We report aggregate CMAC results without releasing the underlying private scripts. Readers seeking access should contact the original authors.

Scope and difficulty. In total the benchmark comprises more than 1,600 lemma-level obligations across these four parts, with the post-quantum developments contributing the largest share. Within every part, obligations span three difficulty levels: "easy" obligations test local proof mechanics, "medium" ones require selecting and following a non-trivial proof strategy, and "hard" ones are the security-proof components—adversary and oracle reasoning, game hops, and probability-bound composition. Because some of the developments were written for older versions of the proof assistant, we port them to current EasyCrypt where feasible and record, per obligation, whether it is machine-checkable today; the obligations that remain are flagged so that the benchmark grows as more developments can be loaded and checked.

C.2. Final Results for MEE-CBC and CMAC

Table 2 summarizes the completed proofs and the remaining obligations in the two case studies. For MEE-CBC, we report the three final theorem targets together with the helper lemmas that the decomposer proposed for them. For CMAC, we report individual lemma and equivalence statements, each counted once even if it was attempted more than once. These counts refer to different proof tasks and are not compared as success rates.

MEE-CBC. The agents receive the security model and selected targets with their intermediate human lemmas removed. Using Opus 5 for decomposition and Opus 4.8 for lemma proving, ShannonProver constructs auxiliary lemmas and their proofs in separate runs for the selected tasks. It completes all three final targets: confidentiality (`RCPA_security`), plaintext integrity (`PTXT_security`), and functional correctness (`mee_correct`). Across the three targets, the decomposer proposed 33 helper lemmas, and the lemma prover proved all 33 together with the three targets; each target run took 23–41 minutes.

The experiment reuses existing verified dependencies, including the CBC library and its security proof. The agent-generated proofs are combined with retained proofs in the final source files. Thus, the result establishes the three target theorems without requiring the agents to regenerate every proof in the supporting libraries. The final sources are available in the MEE-CBC proof browser.

CMAC. For CMAC, we evaluate lemma proving only (Phase III): the lemma and equivalence statements come from the original development, and we report how many of them now have proofs that EasyCrypt accepts. The original development contains 333 named statements, which we kept unchanged except for updates to EasyCrypt syntax. Of these, 299 have agent-generated proofs: 239 were proved directly, and 60 reuse a proof found for a matching statement in another file. Five keep their existing proofs, and 29 remain unfinished, including the final theorem `CMAC_security`. One file, which relates EMAC to CMAC, could not be loaded and referred to missing definitions; we reconstructed it and report its 11 statements separately: nine have agent-generated proofs and two remain unfinished. All proofs are

TABLE 2. FINAL RESULTS FOR MEE-CBC AND CMAC. FOR MEE-CBC, THE DECOMPOSER PROPOSED THE HELPER LEMMAS FOR EACH FINAL THEOREM; FOR CMAC, THE LEMMA AND EQUIVALENCE STATEMENTS COME FROM THE ORIGINAL DEVELOPMENT.

Case study and scope	Completed proofs	Unfinished proofs
MEE-CBC: confidentiality (`RCPA_security`)	Final theorem and its 8 helper lemmas	None for this target
MEE-CBC: plaintext integrity (`PTXT_security`)	Final theorem and its 4 helper lemmas	None for this target
MEE-CBC: functional correctness (`mee_correct`)	Final theorem and its 21 helper lemmas	None for this target
CMAC: 333 original statements, with syntax changes where needed	304: 299 with AI-generated proofs and 5 with existing proofs	29, including `CMAC_security`
CMAC: 11 reconstructed EMAC-to-CMAC statements	9, all with AI-generated proofs	2: `pr_PRP_PRF_ECBC` and `CMAC`

MEE-CBC uses an existing verified CBC library. Of the 299 CMAC statements with AI-generated proofs, 239 have newly found proofs and 60 reuse an AI-generated proof checked for another matching statement. CMAC lemma proofs are checked under the supplied dependencies, some of which remain unproved; the complete CMAC security proof is unfinished.

checked against the supplied dependencies, some of which remain unproved. The counts therefore show how much of the proof work is done, not a complete CMAC security proof; since lemmas differ in difficulty, they also do not measure how much expert effort is saved.

C.3. Expert Comments on MEE-CBC and CMAC

We consulted experts who participated in the developments or had substantial knowledge of them, and received their comments on how LLM agents can accelerate their work.

MEE-CBC. This development started from an existing proof that CBC mode is secure with uniformly random IVs, which was composed with a new proof that MAC-then-Encode-then-Encrypt is INT-PTXT secure. This new proof was formalized by an expert EasyCrypt user with one person-week of effort. A further proof refining the security result to an implementation model (a C-like imperative model in EasyCrypt) took an additional person-week from the same expert. The main source of expert judgment in this development is in choosing the right abstractions and modularizing the proof; once this structure was fixed, much of the remaining work was mechanical proof engineering. Experts commented that the novel part of the MEE-CBC proof involved little human insight, and would have been automated with LLM agents.

CMAC. The proof for CMAC took three person-months by ordinary users of EasyCrypt. This effort also included identifying and resolving small issues in the pen-and-paper proof, as well as making nontrivial design decisions for a part of the argument that could not be formalized as written. (The formalized bound is looser than the pen-and-paper bound as a result.) The experts suggested that the CMAC proof could have been greatly accelerated with LLM agents, not only by reducing the amount of mechanical proof work, but also by surfacing issues in the pen-and-paper proof earlier and helping inform the high-level proof decisions that affect the theorem statement.

Appendix D. Baseline Failure Examples and Standout Agent Proofs

Baseline runs, in which the agent interacts directly with EasyCrypt, show the three failure patterns of §2.

❶ *Unconscious lowering: inlining before applying an equivalence.* The agent recognized that the oracle equivalence `equ_cc` should be applied at a procedure call and recovered its quantified parameters. It then used `inline *`. Although EasyCrypt accepted this tactic, it removed the call structure needed to apply the equivalence. The agent spent 35 minutes on the resulting unsuccessful proof attempt. The baseline interface did not warn that inlining would make the intended lemma inapplicable.

❷ *Semantic mislocalization: choosing a lemma by its name.* The agent searched around `IFinRO`/`IndRO` and chose a lemma whose name looked relevant. It spent more than 25 minutes trying to use `pr_RO_FinRO_D` and inlining eager and lazy loops. The useful next step was to apply `byequiv` and reason about the calls to `chacha_enc2` and `poly_mac2`. The agent followed name similarity without matching the lemma to the procedure calls in the current goal.

❸ *Agent flinch: not applying an identified lemma.* The agent repeatedly identified the right strategy: use `byequiv`, apply the bridge lemma `pr_CCP_OCCP`, and use the PRF distinguishing argument. It proposed this plan four times but never applied the bridge lemma. Instead, it tried local tactics such as `proc; inline *` and `wp; skip; smt()`, or restarted the proof. The difficulty was in supplying the arguments needed to apply the lemma. The baseline interface did not provide the typed argument information needed to turn the plan into a proof step.

Standout agent proofs. Agent proofs are sometimes better than the human proofs. For five ChaChaPoly lemmas, the experts judged the agent proof better in concision, readability, and semantic idiom (Fig. 10). We describe the two marked with stars in the figure.

`chacha_spec`*: stating the loop invariant with the specification.* The lemma states that the loop implementing ChaCha20 encryption computes the functional specification `gen_CTR_encrypt_bytes`, which encrypts the message block by block with counters $1, 2, \ldots$. The difficulty is the loop invariant. The human proof states it in terms of the number of blocks that remain to be encrypted, so each iteration must show by integer arithmetic that this number decreases by exactly one, with a case split on whether a full block remains (23 lines). The agent instead states the invariant with the specification itself: the ciphertext produced so far, followed by the encryption of the remaining message from the current counter, equals the encryption of the whole message. One iteration is then exactly the existing lemma that removes one block from the encryption (`gen_CTR_encrypt_bytes_cons`), and loop exit is the lemma for the empty message (`gen_CTR_encrypt_bytes0`). The arithmetic is still needed, but it is already done inside these lemmas, so the proof takes 7 lines.

`step4_badi`*: describing the bad event exactly.* This lemma is one step of a union bound in the integrity proof. Each call to an internal procedure samples a fresh tag and compares it with a list of earlier tags; the game records every compared pair in a list `lbad1`, and a collision in any pair is a bad event. To bound the probability of a collision, the proof bounds each position i of `lbad1` separately. For this, it uses a modified game that, in the call that fills position i, records in a flag `badi` whether the fresh tag equals the earlier tag at that position. The lemma states that a collision at position i in the original game is no more likely than `badi` in the modified game, and the proof relates the two games one call at a time. The human proof (73 lines) carries an auxiliary invariant relating the length of `lbad1` to the number of queries so far, and distinguishes four cases for each call: no pair is added, or position i was filled earlier, is filled by this call, or is still empty afterwards. The agent's proof (41 lines) states the relation between the two games directly: before position i is filled, `badi` is false, and after it is filled, `badi` holds exactly when the pair at position i collides. Each call then only needs to check whether it fills position i, and the auxiliary invariant on the length of `lbad1` is no longer needed.

Appendix E. Ethics Considerations

Research scope. This work studies automated construction of formal cryptographic proofs. It does not involve a recruited human-subject study, collect or process sensitive personal data, or deploy experiments on third-party systems. It does not report new vulnerabilities in deployed systems requiring coordinated disclosure. The permissions for the cryptographic developments used in our evaluation, including the non-public CMAC development, and the use of generative AI are discussed in Section 9.

Benefits and limitations of proof automation. The intended benefit of ShannonProver is to reduce the effort needed to obtain machine-checked cryptographic proofs. A risk is that users may overinterpret an accepted proof as an unconditional guarantee about a deployed system. The guarantee is limited to the formal statement, its assumptions, and the trusted verification infrastructure; it does not establish that a model captures every property of an implementation. Users must therefore review the specification and theorem statements, and check the final proof scripts with EasyCrypt. An unsuccessful search also does not establish that a theorem is false. These boundaries are essential to using proof automation to support, rather than replace, judgments about cryptographic security.